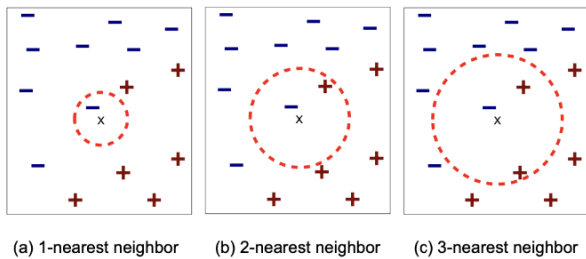


## \* k Nearest-Neighbors Classification

- To classify an example  $x$ , select  $k$  data points of the training set that are most similar to  $x$ .
- Assign the most frequent class among the  $k$  selected.

### >> Instance-based methods

- Simplest form of learning, training dataset is the model itself
- Training dataset is searched for the instances that are more similar to the unlabeled instance
- Similarity function defines what's learned
- Lazy learning: nothing happens until new unlabeled instance must be classified
- Methods: "Rule Learning", "Case Base Reasoning", "k-Nearest Neighbors"



### - How many neighbors?

If  $k$  too small  $\rightarrow$  classification may be sensitive to noise


If  $k$  too large  $\rightarrow$  neighborhood may include dissimilar instances

- What similarity measures?  $\rightarrow$  same ones we applied for clustering =)
- $\rightarrow$  like clustering, we must apply normalization when needed.

### - Basic Approach:


- Linear scan of the data
- Classification time for single distance depends on number of data points and number of variables  $O(nd)$   $\rightarrow$  huge if training set large!
- Nearest Neighbor search can be speed up using: **KD-Trees**, **Ball-Trees**

### - KD-Trees:



- Split the space hierarchically using a tree generated by the data
- Add points iteratively to tree
- Each point fall in a leaf and splits region around it based on value of one of its attributes
- Search for nearest neighbors: Navigate tree to reach leaf and check. Backtrack up tree to check nearby regions. Until all  $k$ -nearest neighbors found
- Search complexity:  $O(\log n)$  for balanced tree
- To build a good tree;
  - Find a good split point and split direction
  - Possible split direction: greatest variance, possible split point: median value along that direction
  - If data is skewed; use value closest to the mean (rather than median)

## - Ball Trees:



- Corners in high-dim space may mean query ball intersects with many regions  $\rightarrow$  limitation for KD-tree
- Can use hyper spheres (balls) instead of hyper rectangles
- Ball tree organizes the data into a tree of  $k$ -dimensional hyperspheres
- Balls may allow for a better fit to the data and thus more efficient search

## - KNN Regression:

- 1) Given a data set  $(x_1, y_1), \dots, (x_n, y_n)$
- 2) Given a query point  $x_q$
- 3) The value  $y_q$  associated to  $x_q$  is computed as a local interpolation of the targets associated to neighbor points
  - Prediction can use average (or weighted average) of  $k$ -nearest targets
  - Prediction can use kernel functions that take the distance as input and return a weight

## - Discussion:

- KNN is often very accurate but slow since it scans entire training data to derive prediction
- Assumes all attributes are equally important  $\rightarrow$  may need attribute selection or weights
- For noisy instances: Majority vote over  $k$ -nearest neighbors  
Remove noisy instances from dataset (difficult = c)